

Exercise 9 - MCMC - Gibbs Sampling, Metropolis-Hastings Sampling

Stefan Merdian

2025-01-12

```
rho <- 0.5  
M <- 30000
```

a) Implement a Gibbs sampler

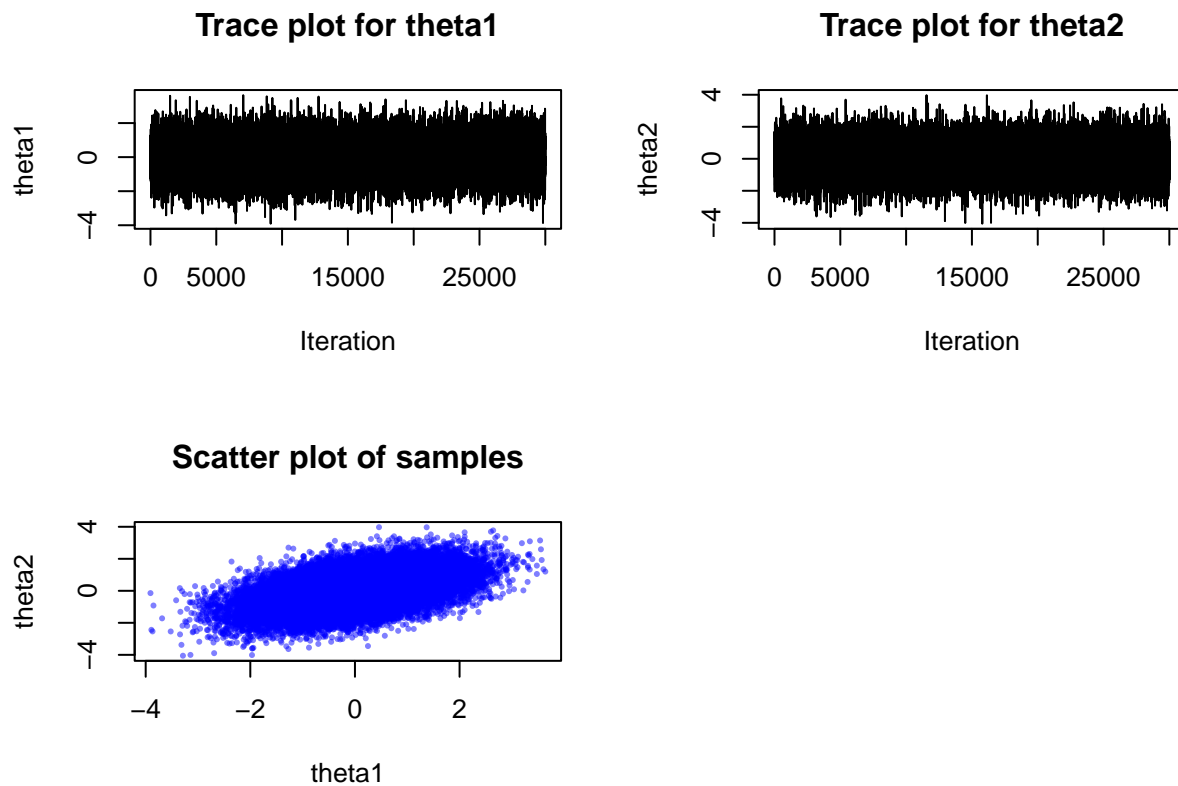
```
gibbs_sampler <- function(rho, M) {  
  samples <- matrix(0, nrow = M, ncol = 2)  
  colnames(samples) <- c("theta1", "theta2")  
  
  theta1 <- 0  
  theta2 <- 0  
  
  for (i in 1:M) {  
    mu1 <- rho * theta2  
    sigma1 <- sqrt(1 - rho^2)  
    theta1 <- rnorm(1, mean = mu1, sd = sigma1)  
  
    mu2 <- rho * theta1  
    sigma2 <- sqrt(1 - rho^2)  
    theta2 <- rnorm(1, mean = mu2, sd = sigma2)  
  
    samples[i, ] <- c(theta1, theta2)  
  }  
  
  return(samples)  
}
```

We use the Gibbs sampler to generate samples from a bivariate normal distribution with parameters ρ and $M=30000$. Then we will show the sampled values of θ_1 and θ_2 over iterations.

Scatter plot: Displays the relationship between θ_1 and θ_2

```
set.seed(123)  
gibbs_samples <- gibbs_sampler(rho, M)  
  
par(mfrow = c(2, 2))  
  
plot(gibbs_samples[, 1], type = "l", main = "Trace plot for theta1", xlab = "Iteration", ylab = "theta1")
```

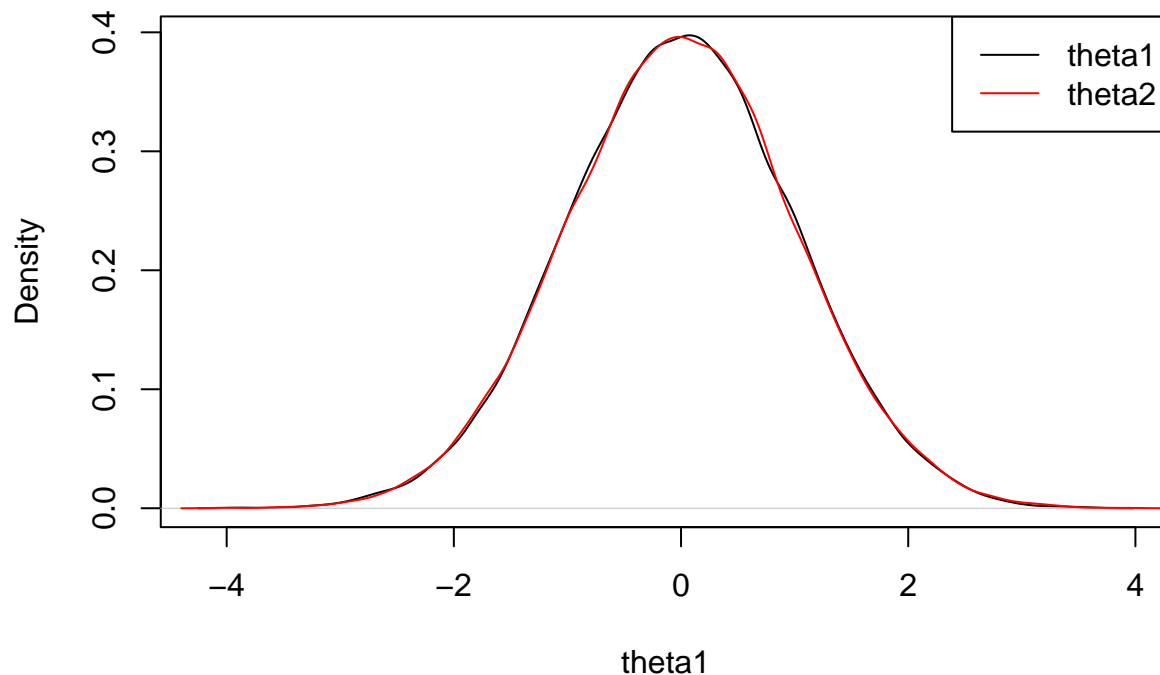
```
plot(gibbs_samples[, 2], type = "l", main = "Trace plot for theta2", xlab = "Iteration", ylab = "theta2")
plot(gibbs_samples, main = "Scatter plot of samples", xlab = "theta1", ylab = "theta2", pch = 16, cex =
```



We also estimate the probability density functions of θ_1 and θ_2 from the Gibbs samples.

```
density_theta1 <- density(gibbs_samples[, 1])
plot(density_theta1, main = "Density plot for theta1", xlab = "theta1", ylab = "Density")
density_theta2 <- density(gibbs_samples[, 2])
lines(density_theta2, col = "red")
legend("topright", legend = c("theta1", "theta2"), col = c("black", "red"), lty = 1)
```

Density plot for theta1



```
# Print summary of the samples
summary(gibbs_samples)
```

```
##      theta1      theta2
## Min.   :-3.908289  Min.   :-4.055195
## 1st Qu.: -0.678665  1st Qu.: -0.672311
## Median :  0.004087  Median :  0.004173
## Mean   :  0.000094  Mean    :  0.001358
## 3rd Qu.:  0.669110  3rd Qu.:  0.669916
## Max.   :  3.638826  Max.    :  3.973645
```

b) Metropolis–Hastings algorithm with block-wise update

1. Initialization: Start with an initial value for $\theta = (0,0)$
2. Target density: Define the bivariate normal target density function
3. Proposal step: Propose a new value for θ_1 or θ_2 using a normal distribution centered at the current value with a standard deviation (`proposal_sd = 1`).
4. Acceptance step: Calculate the acceptance probability for the proposed value based on the ratio of target densities. Accept or reject the proposed value based on a random draw.
5. Repeat: Update θ iteratively for MM iterations, storing the samples at each step.

```
mh_sampler <- function(rho, M) {
  samples <- matrix(0, nrow = M, ncol = 2)
  colnames(samples) <- c("theta1", "theta2")

  theta <- c(0, 0)
```

```

target_density <- function(theta) {
  exp(-0.5 / (1 - rho^2) * (theta[1]^2 - 2 * rho * theta[1] * theta[2] + theta[2]^2))
}

proposal_sd <- 1

for (i in 1:M) {
  for (j in 1:2) {
    theta_proposed <- theta
    theta_proposed[j] <- rnorm(1, mean = theta[j], sd = proposal_sd)

    acceptance_prob <- min(1, target_density(theta_proposed) / target_density(theta))

    if (runif(1) < acceptance_prob) {
      theta <- theta_proposed
    }
  }

  samples[i, ] <- theta
}

return(samples)
}

```

```
mh_samples <- mh_sampler(rho, M)
```

We also estimate the probability density functions of θ_1 and θ_2 from the MH algorithm.

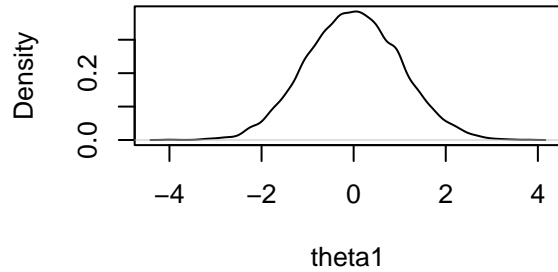
```

par(mfrow = c(2, 2))
plot(density(mh_samples[, 1]), main = "MH: Density plot for theta1", xlab = "theta1", ylab = "Density")
plot(density(mh_samples[, 2]), main = "MH: Density plot for theta2", xlab = "theta2", ylab = "Density")

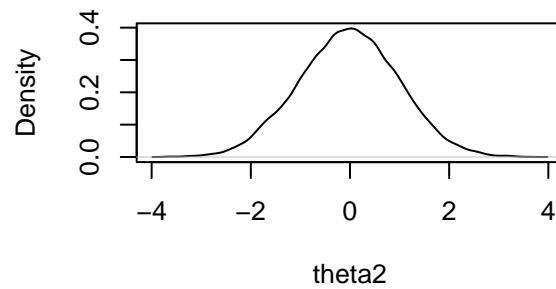
# Scatter plot
plot(mh_samples, main = "MH: Scatter plot of samples", xlab = "theta1", ylab = "theta2", pch = 16, cex = 1)

```

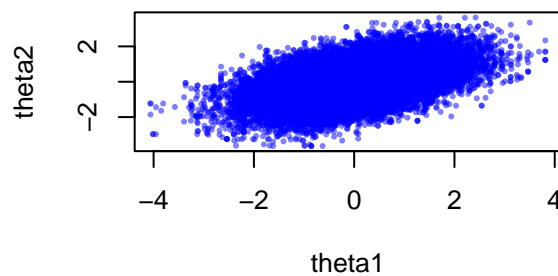
MH: Density plot for theta1



MH: Density plot for theta2



MH: Scatter plot of samples



```
cat("Summary of Metropolis-Hastings samples:\n")
```

```
## Summary of Metropolis-Hastings samples:
```

```
print(summary(mh_samples))
```

```
##      theta1      theta2
##  Min.   :-4.062732  Min.   :-3.63653
##  1st Qu.: -0.699890  1st Qu.: -0.70048
##  Median :-0.007292  Median :-0.01720
##  Mean   :-0.004950  Mean   :-0.02552
##  3rd Qu.: 0.684021  3rd Qu.: 0.65206
##  Max.    : 3.812306  Max.    : 3.63415
```

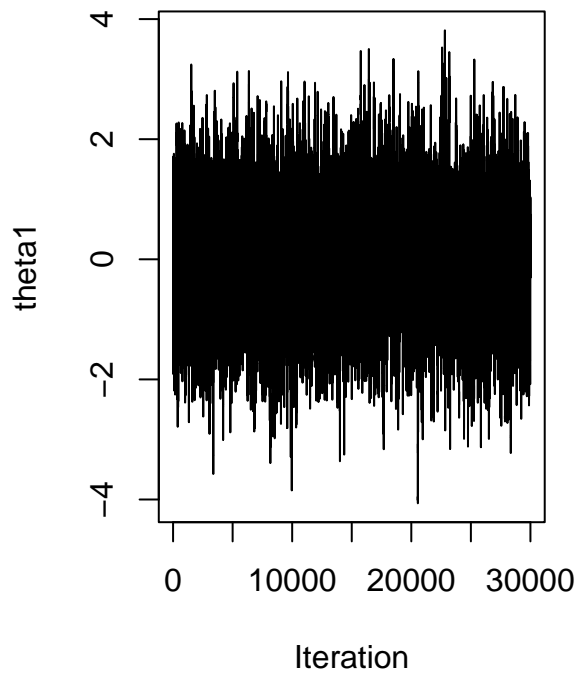
chain diagnostic

```
library(coda)
```

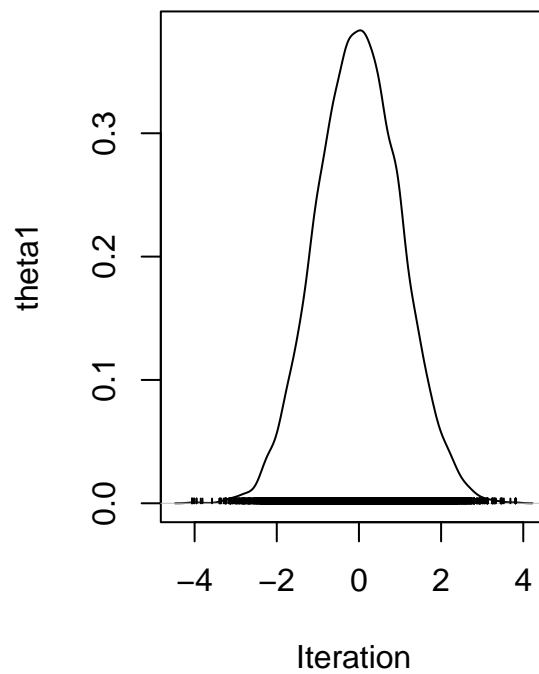
```
## Warning: package 'coda' was built under R version 4.4.2
```

```
gibbs_mcmc <- as.mcmc(gibbs_samples)
mh_mcmc <- as.mcmc(mh_samples)
par(mfrow = c(2, 1))
plot(mh_mcmc[, 1], type = "l", main = "MH: Trace plot for theta1", xlab = "Iteration", ylab = "theta1")
```

MH: Trace plot for theta1

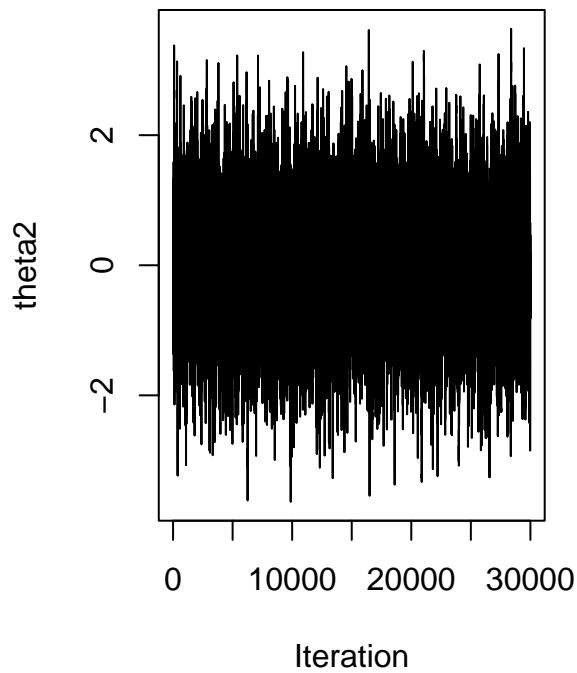


MH: Trace plot for theta1

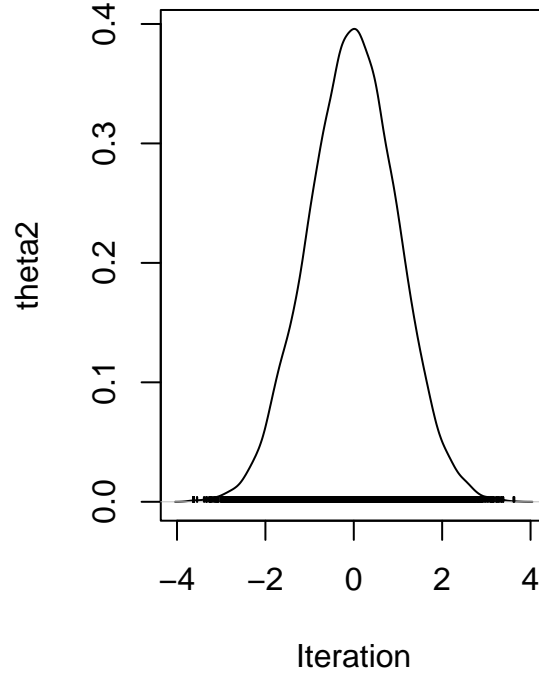


```
plot(mh_mcmc[, 2], type = "l", main = "MH: Trace plot for theta2", xlab = "Iteration", ylab = "theta2")
```

MH: Trace plot for theta2

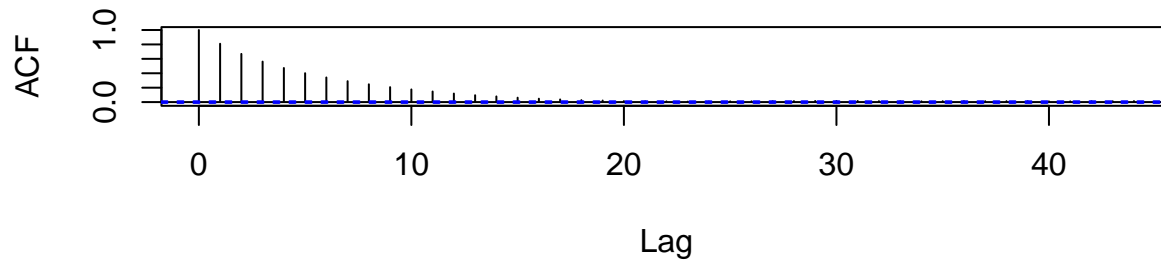


MH: Trace plot for theta2

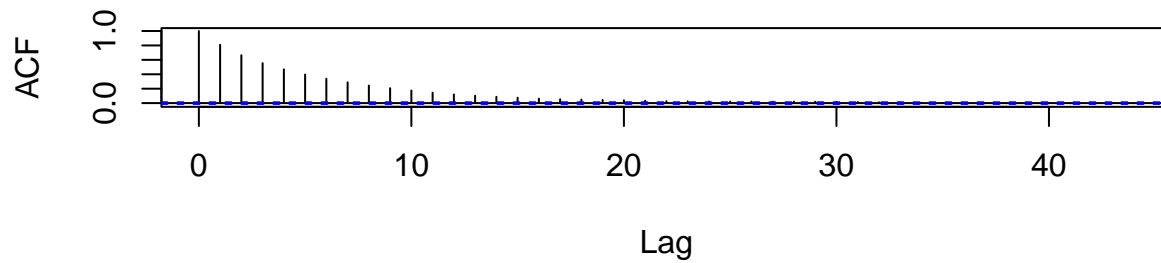


```
acf(mh_mcmc[, 1], main = "MH: Autocorrelation for theta1")  
acf(mh_mcmc[, 2], main = "MH: Autocorrelation for theta2")
```

MH: Autocorrelation for theta1



MH: Autocorrelation for theta2



Trace

Plots:

- The left plot (in 1/2) for both θ_1 and θ_2 , show the sampled values over the iterations
- The chains appear well-mixed, with no visible trends or drifts, indicating that the algorithm is exploring the parameter space effectively and has likely converged.

Density plots (right plot in 1/2):

- Also for both the densities are symmetric and approximately centered around 0, consistent with the bivariate normal distribution.

Autocorrelation Plots:

- The autocorrelation decreases rapidly with lag, which is a good sign of efficient mixing.
- Low autocorrelation indicates that successive samples are less dependent on each other, improving the effective sample size.

```
# Effective sample size
cat("Effective sample size (MH):\n")
```

```
## Effective sample size (MH):
```

```
print(effectiveSize(mh_mcmc))
```

```
##   theta1   theta2
## 2748.768 2728.083
```

```
# Geweke diagnostic
cat("Geweke diagnostic (MH):\n")
```

```
## Geweke diagnostic (MH):
```

```
print(geweke.diag(mh_mcmc))
```

```
##
## Fraction in 1st window = 0.1
## Fraction in 2nd window = 0.5
##
## theta1 theta2
## 0.3385 0.8182
```

Effective Sample Size:

- The effective sample size indicates the number of independent samples the chain effectively provides.
- While the chain has 30,000 iterations, the effective sample size is smaller due to autocorrelation. However, these values are high, showing the chain mixes well.

Geweke Diagnostic:

- These values represent z-scores. Values near 0 indicate no significant difference between the two segments, suggesting convergence.
- Both 1 and 2 have values close to 0, confirming the chains have likely converged.